

17.1 Suppose that there is a database system that never fails. Is a recovery manager required for this system?

Even in this case the recovery manager is needed to perform roll-back of aborted transactions.

17.3 Database-system implementers have paid much more attention to the ACID properties than have file-system implementers. Why might this be the case?

Database systems usually perform crucial tasks whose effects need to be atomic and durable, and whose outcome affects the real world in a permanent manner. Examples of such tasks are monetary transactions, seat bookings etc. Hence the ACID properties have to be ensured.

17.15 Consider the following two transactions:

T13:

```
read(A);
read(B);
if A = 0 then B := B + 1;
write(B).
```

T14:

```
read(B);
read(A);
if B = 0 then A := A + 1;
write(A).
```

Let the consistency requirement be $A = 0 \vee B = 0$, with $A = B = 0$ as the initial values.

a. Show that every serial execution involving these two transactions preserves the consistency of the database.

Case 1: $T13 \rightarrow T14$

1. Initial state: $A = 0, B = 0$.
2. T13:
 - Reads $A = 0$, then reads $B = 0$.
 - Since $A = 0, B := B + 1 = 1$.
 - Writes $B = 1$.
 - State after T13: $A = 0, B = 1$.
3. T14:
 - Reads $B = 1$, then reads $A = 0$.
 - Since $B \neq 0$, no changes to A .
 - Writes $A = 0$.
 - State after T14: $A = 0, B = 1$.
4. Consistency Check: $A = 0 \vee B = 0$ is satisfied ($A = 0$).

Case 2: $T14 \rightarrow T13$

1. Initial state: $A = 0, B = 0$.
2. T14:
 - Reads $B = 0$, then reads $A = 0$.
 - Since $B = 0, A := A + 1 = 1$.
 - Writes $A = 1$.
 - State after T14: $A = 1, B = 0$.
3. T13:
 - Reads $A = 1$, then reads $B = 0$.
 - Since $A \neq 0$, no changes to B .
 - Writes $B = 0$.
 - State after T13: $A = 1, B = 0$.
4. Consistency Check: $A = 0 \vee B = 0$ is satisfied ($B = 0$).

Conclusion

In both serial executions, the consistency requirement $A=0 \vee B=0$ is preserved. This shows that every serial execution involving these transactions maintains the consistency of the database.

b. Show a concurrent execution of T13 and T14 that produces a nonserializable schedule.

Let us consider the operations in T_{13} and T_{14} :

- T_{13} : `read(A), read(B), if A = 0 then B := B + 1, write(B)`
- T_{14} : `read(B), read(A), if B = 0 then A := A + 1, write(A)`

Concurrent Schedule

We interleave the operations of T_{13} and T_{14} as follows:

1. T_{13} : `read(A)` → Reads $A = 0$.
2. T_{14} : `read(B)` → Reads $B = 0$.
3. T_{13} : `read(B)` → Reads $B = 0$.
4. T_{14} : `read(A)` → Reads $A = 0$.
5. T_{13} : Since $A = 0$, sets $B := B + 1 = 1$.
6. T_{14} : Since $B = 0$, sets $A := A + 1 = 1$.
7. T_{13} : `write(B)` → Writes $B = 1$.
8. T_{14} : `write(A)` → Writes $A = 1$.

Resulting State

- After the schedule completes: $A = 1, B = 1$.

Consistency Check

The consistency requirement is $A = 0 \vee B = 0$. However, the final state $A = 1, B = 1$ violates the consistency condition, as neither $A = 0$ nor $B = 0$ holds.

Nonserializability

This schedule is **nonserializable** because:

- In both possible serial schedules ($T_{13} \rightarrow T_{14}$ or $T_{14} \rightarrow T_{13}$), at least one of A or B would remain 0, ensuring the consistency condition.
- The interleaving above introduces a state ($A = 1, B = 1$) that cannot be obtained from any serial schedule.

c. Is there a concurrent execution of T13 and T14 that produces a serializable schedule?

For part (c), the answer is **no**, there is no concurrent execution of T13 and T14 that produces a **serializable schedule**. There is no concurrent execution of T13 and T14 that produces a **prserializable schedule**.

17.16 Give an example of a serializable schedule with two transactions such that the order in which the transactions commit is different from the serialization order.

Transactions

We define two transactions, T_1 and T_2 , as follows:

- T_1 :
`read(X) → write(X) → commit.`
- T_2 :
`read(X) → write(X) → commit.`

Serializable Schedule

A **schedule** interleaving the operations of T_1 and T_2 is:

1. T_1 : `read(X)`
2. T_2 : `read(X)`
3. T_2 : `write(X)`
4. T_2 : `commit`
5. T_1 : `write(X)`
6. T_1 : `commit.`

Serialization Order

The **serialization order** of this schedule is $T2 \rightarrow T1$:

- $T2$'s `write(X)` logically occurs before $T1$'s `write(X)`.
- The final state of X is determined by $T1$, which overwrites $T2$'s changes.

Commit Order

The **commit order** is $T1 \rightarrow T2$:

- $T1$ commits after $T2$ finishes its `write(X)` but before $T2$ commits.

Why is this Serializable?

The schedule is serializable because:

1. The operations can be reordered into a valid serial execution ($T2 \rightarrow T1$) without affecting the final database state.
2. There is no cyclic dependency in the precedence graph of the transactions.